

# CPSC 330

# Applied Machine Learning

## Lecture 10: Regression Evaluation Metrics

UBC 2022-23

Instructor: Mathias Lécuyer

### Announcement

- Midterm next week, on Wednesday, Feb. 15, from 7:00pm to 8:15pm:
  - CPSC 330 Section 202 students (you) will write the exam in ESB 1012 (<https://learningspaces.ubc.ca/classrooms/esb-1012> (<https://learningspaces.ubc.ca/classrooms/esb-1012>)):
  - More details on piazza (<https://piazza.com/class/lcgo6c2ncl06el/post/283>): <https://piazza.com/class/lcgo6c2ncl06el/post/283> (<https://piazza.com/class/lcgo6c2ncl06el/post/283>).
- There is a [piazza poll](https://piazza.com/class/lcgo6c2ncl06el/post/316) (<https://piazza.com/class/lcgo6c2ncl06el/post/316>) for topics to cover in the review session (next Tuesday, Feb 14): <https://piazza.com/class/lcgo6c2ncl06el/post/316> (<https://piazza.com/class/lcgo6c2ncl06el/post/316>).

## Imports

```
In [41]: 1 import matplotlib.pyplot as plt
2 import numpy as np
3 import pandas as pd
4 from sklearn.compose import (
5     ColumnTransformer,
6     TransformedTargetRegressor,
7     make_column_transformer,
8 )
9 from sklearn.dummy import DummyRegressor
10 from sklearn.ensemble import RandomForestRegressor
11 from sklearn.impute import SimpleImputer
12 from sklearn.linear_model import LinearRegression, Ridge, RidgeCV
13 from sklearn.metrics import make_scorer, mean_squared_error, r2_score
14 from sklearn.model_selection import cross_val_score, cross_validate, tr
15 from sklearn.pipeline import Pipeline, make_pipeline
16 from sklearn.preprocessing import OneHotEncoder, OrdinalEncoder, Standa
17 from sklearn.tree import DecisionTreeRegressor
18
19 %matplotlib inline
```

```
In [42]: 1 import warnings
2
3 warnings.simplefilter(action="ignore", category=FutureWarning)
```

## Some clarifications on last lecture

- Formula for false positive rate: Fraction of false positives out of all negative examples:

$$FPR = \frac{FP}{FP + TN} = \frac{FP}{N}$$

- This was in context of the ROC curve, which is TPR (a.k.a recall, a.k.a sensitivity) as function of FPR:

$$TPR = \frac{TP}{TP + FN} = \frac{TP}{P}, \quad FPR = \frac{FP}{FP + TN}$$

- Recall is also called sensitivity:

$$\text{Recall} = \text{Sensitivity} = \frac{TP}{TP + FN} = \frac{TP}{P}$$

- Note that TPR/recall/sensitivity is about the positive class (how much of it we find), while FPR is really about the negative class (how much of it we mispredict).
- Precision, recall, f1 score, are only about the positive label:

$$\text{Precision} = \frac{TP}{TP + FP}, \quad \text{Recall} = \text{Sensitivity} = \frac{TP}{TP + FN} = \frac{TP}{P}$$

Processing math: 100%

- The positive class is assumed to be the class label 1 by default. This is configurable through the `pos_label` parameter.

## Learning outcomes

From this lecture, students are expected to be able to:

- Carry out feature transformations on somewhat complicated dataset.
- Visualize transformed features as a dataframe.
- Use `Ridge` and `RidgeCV`.
- Explain how `alpha` hyperparameter of `Ridge` relates to the fundamental tradeoff.
- Examine coefficients of transformed features.
- Appropriately select a scoring metric given a regression problem.
- Interpret and communicate the meanings of different scoring metrics on regression problems.
  - `MSE`, `RMSE`,  $R^2$ , `MAPE`
- Apply log-transform on the target values in a regression problem with `TransformedTargetRegressor`.

## Dataset

In this lecture, we'll be using [Kaggle House Prices dataset \(https://www.kaggle.com/c/home-data-for-ml-course/\)](https://www.kaggle.com/c/home-data-for-ml-course/). As usual, to run this notebook you'll need to download the data. For this dataset, train and test have already been separated. We'll be working with the train portion in this lecture.

```
In [43]: 1 df = pd.read_csv("../data/housing-kaggle/train.csv")
          2 train_df, test_df = train_test_split(df, test_size=0.10, random_state=1)
          3 train_df.head()
```

```
Out[43]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	L
302	303	20	RL	118.0	13704	Pave	NaN	IR1		Lvl
767	768	50	RL	75.0	12508	Pave	NaN	IR1		Lvl
429	430	20	RL	130.0	11457	Pave	NaN	IR1		Lvl
1139	1140	30	RL	98.0	8731	Pave	NaN	IR1		Lvl
558	559	60	RL	57.0	21872	Pave	NaN	IR2		HLS

5 rows × 81 columns

- The supervised machine learning problem is predicting housing price given features associated with properties.
- Here, the target is `SalePrice`, which is continuous. So it's a **regression problem** (as opposed to classification).

```
In [44]: 1 train_df.shape
```

```
Out[44]: (1314, 81)
```

## Let's separate x and y

```
In [45]: 1 x_train = train_df.drop(columns=["SalePrice"])
        2 y_train = train_df["SalePrice"]
        3
        4 x_test = test_df.drop(columns=["SalePrice"])
        5 y_test = test_df["SalePrice"]
```

## EDA

```
In [46]: 1 train_df.describe()
```

```
Out[46]:
```

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBui
<b>count</b>	1314.000000	1314.000000	1089.000000	1314.000000	1314.000000	1314.000000	1314.000000
<b>mean</b>	734.182648	56.472603	69.641873	10273.261035	6.076104	5.570015	1970.99543
<b>std</b>	422.224662	42.036646	23.031794	8997.895541	1.392612	1.112848	30.19812
<b>min</b>	1.000000	20.000000	21.000000	1300.000000	1.000000	1.000000	1872.00000
<b>25%</b>	369.250000	20.000000	59.000000	7500.000000	5.000000	5.000000	1953.00000
<b>50%</b>	735.500000	50.000000	69.000000	9391.000000	6.000000	5.000000	1972.00000
<b>75%</b>	1099.750000	70.000000	80.000000	11509.000000	7.000000	6.000000	2000.00000
<b>max</b>	1460.000000	190.000000	313.000000	215245.000000	10.000000	9.000000	2010.00000

8 rows × 38 columns

```
In [47]: 1 train_df.info()
```

Processing math: 100%

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 1314 entries, 302 to 1389
```

```
Data columns (total 81 columns):
```

#	Column	Non-Null Count	Dtype
0	Id	1314 non-null	int64
1	MSSubClass	1314 non-null	int64
2	MSZoning	1314 non-null	object
3	LotFrontage	1089 non-null	float64
4	LotArea	1314 non-null	int64
5	Street	1314 non-null	object
6	Alley	81 non-null	object
7	LotShape	1314 non-null	object
8	LandContour	1314 non-null	object
9	Utilities	1314 non-null	object
10	LotConfig	1314 non-null	object
11	LandSlope	1314 non-null	object
12	Neighborhood	1314 non-null	object
13	Condition1	1314 non-null	object
14	Condition2	1314 non-null	object
15	BldgType	1314 non-null	object
16	HouseStyle	1314 non-null	object
17	OverallQual	1314 non-null	int64
18	OverallCond	1314 non-null	int64
19	YearBuilt	1314 non-null	int64
20	YearRemodAdd	1314 non-null	int64
21	RoofStyle	1314 non-null	object
22	RoofMatl	1314 non-null	object
23	Exterior1st	1314 non-null	object
24	Exterior2nd	1314 non-null	object
25	MasVnrType	1307 non-null	object
26	MasVnrArea	1307 non-null	float64
27	ExterQual	1314 non-null	object
28	ExterCond	1314 non-null	object
29	Foundation	1314 non-null	object
30	BsmtQual	1280 non-null	object
31	BsmtCond	1280 non-null	object
32	BsmtExposure	1279 non-null	object
33	BsmtFinType1	1280 non-null	object
34	BsmtFinSF1	1314 non-null	int64
35	BsmtFinType2	1280 non-null	object
36	BsmtFinSF2	1314 non-null	int64
37	BsmtUnfSF	1314 non-null	int64
38	TotalBsmtSF	1314 non-null	int64
39	Heating	1314 non-null	object
40	HeatingQC	1314 non-null	object
41	CentralAir	1314 non-null	object
42	Electrical	1313 non-null	object
43	1stFlrSF	1314 non-null	int64
44	2ndFlrSF	1314 non-null	int64
45	LowQualFinSF	1314 non-null	int64
46	GrLivArea	1314 non-null	int64
47	BsmtFullBath	1314 non-null	int64
48	BsmtHalfBath	1314 non-null	int64
49	FullBath	1314 non-null	int64
50	HalfBath	1314 non-null	int64
51	BedroomAbvGr	1314 non-null	int64

Processing math: 100%

```

52 KitchenAbvGr 1314 non-null int64
53 KitchenQual 1314 non-null object
54 TotRmsAbvGrd 1314 non-null int64
55 Functional 1314 non-null object
56 Fireplaces 1314 non-null int64
57 FireplaceQu 687 non-null object
58 GarageType 1241 non-null object
59 GarageYrBlt 1241 non-null float64
60 GarageFinish 1241 non-null object
61 GarageCars 1314 non-null int64
62 GarageArea 1314 non-null int64
63 GarageQual 1241 non-null object
64 GarageCond 1241 non-null object
65 PavedDrive 1314 non-null object
66 WoodDeckSF 1314 non-null int64
67 OpenPorchSF 1314 non-null int64
68 EnclosedPorch 1314 non-null int64
69 3SsnPorch 1314 non-null int64
70 ScreenPorch 1314 non-null int64
71 PoolArea 1314 non-null int64
72 PoolQC 7 non-null object
73 Fence 259 non-null object
74 MiscFeature 50 non-null object
75 MiscVal 1314 non-null int64
76 MoSold 1314 non-null int64
77 YrSold 1314 non-null int64
78 SaleType 1314 non-null object
79 SaleCondition 1314 non-null object
80 SalePrice 1314 non-null int64
dtypes: float64(3), int64(35), object(43)
memory usage: 841.8+ KB

```

## pandas\_profiler

We do not have `pandas_profiling` in our course environment. You will have to install it in the environment on your own if you want to run the code below.

```
conda install -c conda-forge pandas-profiling
```

```

In [48]: 1 from pandas_profiling import ProfileReport
          2
          3 #profile = ProfileReport(train_df, title="Pandas Profiling Report") #
          4 #profile.to_notebook_iframe()

```

## Feature types

- Do not blindly trust all the info given to you by automated tools.
- How does pandas profiling figure out the data type?
  - You can look at the Python data type and say floats are numeric, strings are categorical.
  - However, in doing so you would miss out on various subtleties such as some of the string features being ordinal rather than truly categorical.
  - Also, it will think free text is categorical.

- In addition to tools such as above, it's important to go through data description to understand the data.
- The data description for our dataset is available [here \(https://www.kaggle.com/c/home-data-for-ml-course/data?select=data\\_description.txt\)](https://www.kaggle.com/c/home-data-for-ml-course/data?select=data_description.txt).

## Feature types

- We have mixed feature types and a bunch of missing values.
- Now, let's identify feature types and transformations.

- Let's get the numeric-looking columns.

```
In [49]: 1 numeric_looking_columns = X_train.select_dtypes(include=np.number).columns
          2 print(numeric_looking_columns)
```

```
['Id', 'MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces', 'GarageYrBlt', 'GarageCars', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal', 'MoSold', 'YrSold']
```

Not all numeric looking columns are necessarily numeric.

```
In [50]: 1 train_df["MSSubClass"].unique()
```

```
Out[50]: array([ 20,  50,  30,  60, 160,  85,  90, 120, 180,  80,  70,  75, 190,
                45,  40])
```

MSSubClass: Identifies the type of dwelling involved in the sale.



```
20 1-STORY 1946 & NEWER ALL STYLES
30 1-STORY 1945 & OLDER
40 1-STORY W/FINISHED ATTIC ALL AGES
45 1-1/2 STORY - UNFINISHED ALL AGES
50 1-1/2 STORY FINISHED ALL AGES
60 2-STORY 1946 & NEWER
70 2-STORY 1945 & OLDER
75 2-1/2 STORY ALL AGES
80 SPLIT OR MULTI-LEVEL
85 SPLIT FOYER
90 DUPLEX - ALL STYLES AND AGES
120 1-STORY PUD (Planned Unit Development) - 1946 & NEWER
```

Also, month sold is more of a categorical feature than a numeric feature.

```
In [51]: 1 train_df["MoSold"].unique() # Month Sold
```

```
Out[51]: array([ 1,  7,  3,  5,  8, 10,  6,  9, 12,  2,  4, 11])
```

```

In [52]: 1 drop_features = ["Id"]
          2 numeric_features = [
          3     "BedroomAbvGr",
          4     "KitchenAbvGr",
          5     "LotFrontage",
          6     "LotArea",
          7     "OverallQual",
          8     "OverallCond",
          9     "YearBuilt",
         10     "YearRemodAdd",
         11     "MasVnrArea",
         12     "BsmtFinSF1",
         13     "BsmtFinSF2",
         14     "BsmtUnfSF",
         15     "TotalBsmtSF",
         16     "1stFlrSF",
         17     "2ndFlrSF",
         18     "LowQualFinSF",
         19     "GrLivArea",
         20     "BsmtFullBath",
         21     "BsmtHalfBath",
         22     "FullBath",
         23     "HalfBath",
         24     "TotRmsAbvGrd",
         25     "Fireplaces",
         26     "GarageYrBlt",
         27     "GarageCars",
         28     "GarageArea",
         29     "WoodDeckSF",
         30     "OpenPorchSF",
         31     "EnclosedPorch",
         32     "3SsnPorch",
         33     "ScreenPorch",
         34     "PoolArea",
         35     "MiscVal",
         36     "YrSold",
         37 ]

```

I've not looked at all the features carefully. It might be appropriate to apply some other encoding on some of the numeric features above.

```

In [53]: 1 set(numeric_looking_columns) - set(numeric_features) - set(drop_features)

```

```

Out[53]: {'MSSubClass', 'MoSold'}

```

We'll treat the above numeric-looking features as categorical features.

- There are a bunch of ordinal features in this dataset.
- Ordinal features with the same scale
  - Poor (Po), Fair (Fa), Typical (TA), Good (Gd), Excellent (Ex)

Processing math: 100%

- These we'll be calling `ordinal_features_reg`.
- Ordinal features with different scales
  - These we'll be calling `ordinal_features_oth`.

```
In [54]: 1 ordinal_features_reg = [
2         "ExterQual",
3         "ExterCond",
4         "BsmtQual",
5         "BsmtCond",
6         "HeatingQC",
7         "KitchenQual",
8         "FireplaceQu",
9         "GarageQual",
10        "GarageCond",
11        "PoolQC",
12    ]
13    ordering = [
14        "Po",
15        "Fa",
16        "TA",
17        "Gd",
18        "Ex",
19    ] # if N/A it will just impute something, per below
20    ordering_ordinal_reg = [ordering] * len(ordinal_features_reg)
21    ordering_ordinal_reg
```

```
Out[54]: [['Po', 'Fa', 'TA', 'Gd', 'Ex'],
['Po', 'Fa', 'TA', 'Gd', 'Ex'],
['Po', 'Fa', 'TA', 'Gd', 'Ex'],
['Po', 'Fa', 'TA', 'Gd', 'Ex'],
['Po', 'Fa', 'TA', 'Gd', 'Ex'],
['Po', 'Fa', 'TA', 'Gd', 'Ex'],
['Po', 'Fa', 'TA', 'Gd', 'Ex'],
['Po', 'Fa', 'TA', 'Gd', 'Ex'],
['Po', 'Fa', 'TA', 'Gd', 'Ex'],
['Po', 'Fa', 'TA', 'Gd', 'Ex']]
```

We'll pass the above as categories in our `OrdinalEncoder`.

- There are a bunch more ordinal features using different scales.
  - These we'll be calling `ordinal_features_oth`.
  - We are encoding them separately.

```
In [55]: 1 ordinal_features_oth = [  
2     "BsmtExposure",  
3     "BsmtFinType1",  
4     "BsmtFinType2",  
5     "Functional",  
6     "Fence",  
7 ]  
8 ordering_ordinal_oth = [  
9     ['NA', 'No', 'Mn', 'Av', 'Gd'],  
10    ['NA', 'Unf', 'LwQ', 'Rec', 'BLQ', 'ALQ', 'GLQ'],  
11    ['NA', 'Unf', 'LwQ', 'Rec', 'BLQ', 'ALQ', 'GLQ'],  
12    ['Sal', 'Sev', 'Maj2', 'Maj1', 'Mod', 'Min2', 'Min1', 'Typ'],  
13    ['NA', 'MnWw', 'GdWo', 'MnPrv', 'GdPrv']  
14 ]
```

The remaining features are categorical features.

```
In [56]: 1 categorical_features = list(  
2         set(X_train.columns)  
3         - set(numeric_features)  
4         - set(ordinal_features_reg)  
5         - set(ordinal_features_oth)  
6         - set(drop_features)  
7     )  
8     categorical_features
```

```
Out[56]: ['MasVnrType',  
          'Neighborhood',  
          'Condition2',  
          'Alley',  
          'SaleCondition',  
          'Electrical',  
          'HouseStyle',  
          'GarageFinish',  
          'RoofStyle',  
          'MoSold',  
          'Exterior1st',  
          'Street',  
          'MSSubClass',  
          'LotShape',  
          'PavedDrive',  
          'Heating',  
          'SaleType',  
          'Utilities',  
          'GarageType',  
          'BldgType',  
          'MiscFeature',  
          'LotConfig',  
          'CentralAir',  
          'LandSlope',  
          'Condition1',  
          'Exterior2nd',  
          'MSZoning',  
          'LandContour',  
          'RoofMatl',  
          'Foundation']
```

- We are not doing it here but we can engineer our own features too.
- Would price per square foot be a good feature to add in here?

## Applying feature transformations

- Since we have mixed feature types, let's use `ColumnTransformer` to apply different transformations on different features types.

```
In [57]: 1 from sklearn.compose import ColumnTransformer, make_column_transformer
2
3 numeric_transformer = make_pipeline(SimpleImputer(strategy="median"), S
4 ordinal_transformer_reg = make_pipeline(
5     SimpleImputer(strategy="most_frequent"),
6     OrdinalEncoder(categories=ordering_ordinal_reg),
7 )
8
9 ordinal_transformer_oth = make_pipeline(
10     SimpleImputer(strategy="most_frequent"),
11     OrdinalEncoder(categories=ordering_ordinal_oth),
12 )
13
14 categorical_transformer = make_pipeline(
15     SimpleImputer(strategy="constant", fill_value="missing"),
16     OneHotEncoder(handle_unknown="ignore", sparse=False),
17 )
18
19 preprocessor = make_column_transformer(
20     ("drop", drop_features),
21     (numeric_transformer, numeric_features),
22     (ordinal_transformer_reg, ordinal_features_reg),
23     (ordinal_transformer_oth, ordinal_features_oth),
24     (categorical_transformer, categorical_features),
25 )
```

## Examining the preprocessed data

```
In [58]: 1 preprocessor.fit(X_train) # Calling fit to examine all the transformers  
        2 preprocessor.named_transformers_
```

```

Out[58]: {'drop': 'drop',
          'pipeline-1': Pipeline(steps=[('simpleimputer', SimpleImputer(strategy
='median')),
                                         ('standardscaler', StandardScaler())]),
          'pipeline-2': Pipeline(steps=[('simpleimputer', SimpleImputer(strategy
='most_frequent')),
                                         ('ordinalencoder',
                                          OrdinalEncoder(categories=[['Po', 'Fa', 'TA', 'Gd', 'E
x'],
                                                                    ['Po', 'Fa', 'TA', 'Gd', 'E
x'],
                                                                    ['Po', 'Fa', 'TA', 'Gd', 'E
x'],
                                                                    ['Po', 'Fa', 'TA', 'Gd', 'E
x'],
                                                                    ['Po', 'Fa', 'TA', 'Gd', 'E
x'],
                                                                    ['Po', 'Fa', 'TA', 'Gd', 'E
x'],
                                                                    ['Po', 'Fa', 'TA', 'Gd', 'E
x'],
                                                                    ['Po', 'Fa', 'TA', 'Gd', 'E
x'],
                                                                    ['Po', 'Fa', 'TA', 'Gd', 'E
x'],
                                                                    ['Po', 'Fa', 'TA', 'Gd', 'E
x'],
                                                                    ['Po', 'Fa', 'TA', 'Gd', 'E
x']])))],
          'pipeline-3': Pipeline(steps=[('simpleimputer', SimpleImputer(strategy
='most_frequent')),
                                         ('ordinalencoder',
                                          OrdinalEncoder(categories=[['NA', 'No', 'Mn', 'Av', 'G
d'],
                                                                    ['NA', 'Unf', 'LwQ', 'Rec',
                                                                    'ALQ', 'GLQ'],
                                                                    ['NA', 'Unf', 'LwQ', 'Rec',
                                                                    'ALQ', 'GLQ'],
                                                                    ['Sal', 'Sev', 'Maj2', 'Maj
1',
                                                                    'Mod', 'Min2', 'Min1', 'Ty
p'],
                                                                    ['NA', 'MnWw', 'GdWo', 'MnPr
v',
                                                                    'GdPrv']])))],
          'pipeline-4': Pipeline(steps=[('simpleimputer',
                                          SimpleImputer(fill_value='missing', strategy='constan
t')),
                                         ('onehotencoder',
                                          OneHotEncoder(handle_unknown='ignore', sparse=Fals
e)))]})

```



```
In [59]: 1 ohe_columns = list(
2         preprocessor.named_transformers_[ "pipeline-4" ]
3         .named_steps[ "onehotencoder" ]
4         .get_feature_names(categorical_features)
5     )
6     new_columns = numeric_features + ordinal_features_reg + ordinal_features
```

```
In [60]: 1 X_train_enc = pd.DataFrame(
2         preprocessor.transform(X_train), index=X_train.index, columns=new_columns
3     )
4     X_train_enc.head()
```

```
Out[60]:
```

	BedroomAbvGr	KitchenAbvGr	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemain
302	0.154795	-0.222647	2.312501	0.381428	0.663680	-0.512408	0.993969	-0.000000
767	1.372763	-0.222647	0.260890	0.248457	-0.054669	1.285467	-1.026793	-0.000000
429	0.154795	-0.222647	2.885044	0.131607	-0.054669	-0.512408	0.563314	-0.000000
1139	0.154795	-0.222647	1.358264	-0.171468	-0.773017	-0.512408	-1.689338	-0.000000
558	0.154795	-0.222647	-0.597924	1.289541	0.663680	-0.512408	0.828332	-0.000000

5 rows × 263 columns

```
In [61]: 1 X_train.shape
```

```
Out[61]: (1314, 80)
```

```
In [62]: 1 X_train_enc.shape
```

```
Out[62]: (1314, 263)
```

We went from 80 features to 263 features!!

## Other possible preprocessing?

- There is a lot of room for improvement ...
- We're just using `SimpleImputer` .
  - In reality we'd want to go through this more carefully.
  - We may also want to drop some columns that are almost entirely missing.
- We could also check for outliers, and do other exploratory data analysis (EDA).
- But for now this is good enough ...

## Model building

### DummyRegressor

```
In [63]: 1 dummy = DummyRegressor()
          2 pd.DataFrame(cross_validate(dummy, X_train, y_train, cv=10, return_train_score=True))
```

```
Out[63]:
```

	fit_time	score_time	test_score	train_score
0	0.001767	0.000557	-0.003547	0.0
1	0.001569	0.000441	-0.001266	0.0
2	0.001287	0.000450	-0.011767	0.0
3	0.002081	0.000792	-0.006744	0.0
4	0.001895	0.000488	-0.076533	0.0
5	0.001080	0.000369	-0.003133	0.0
6	0.001010	0.000363	-0.000397	0.0
7	0.000984	0.000360	-0.003785	0.0
8	0.001093	0.000494	-0.001740	0.0
9	0.001242	0.000619	-0.000117	0.0

### Apply Ridge

- Recall that we are going to use `Ridge()` instead of `LinearRegression()` in this course.
  - It has a hyperparameter `alpha` which controls the fundamental tradeoff.

```
In [64]: 1 lr_pipe = make_pipeline(preprocessor, Ridge())
          2 pd.DataFrame(cross_validate(lr_pipe, X_train, y_train, cv=10, return_train_score=True))
```

```
Out[64]:
```

	fit_time	score_time	test_score	train_score
0	0.067918	0.020969	0.861355	0.911906
1	0.073782	0.020503	0.812301	0.913861
2	0.075481	0.021917	0.775283	0.915963
3	0.077557	0.021604	0.874519	0.910849
4	0.074957	0.021669	0.851969	0.911622
5	0.074832	0.024669	0.826198	0.910176
6	0.074639	0.022231	0.825533	0.913781
7	0.074474	0.021569	0.872238	0.910071
8	0.077648	0.024597	0.196663	0.921448
9	0.072799	0.023196	0.890474	0.908221

Processing math: 100%


- Quite a bit of variation in the test scores.
- Performing poorly in fold 8. Not sure why.

## Tuning the `alpha` hyperparameter of Ridge

- Recall that Ridge has a hyperparameter `alpha` that controls the fundamental tradeoff.
- This is like `C` in `LogisticRegression` but, annoyingly, `alpha` is the opposite of `C`: large `C` is like small `alpha` and vice versa.
- Smaller `alpha`: more complex model, more variance, lower training error (overfitting)

```
In [65]: 1 alphas = 10.0 ** np.arange(-5, 5, 1)
2 train_scores = []
3 cv_scores = []
4 for alpha in alphas:
5     lr = make_pipeline(preprocessor, Ridge(alpha=alpha))
6     results = cross_validate(lr, X_train, y_train, return_train_score=True)
7     train_scores.append(np.mean(results["train_score"]))
8     cv_scores.append(np.mean(results["test_score"]))
```

```
In [66]: 1 plt.semilogx(alphas, train_scores, label="train")
2 plt.semilogx(alphas, cv_scores, label="cv")
3 plt.legend()
4 plt.xlabel("alpha")
5 plt.ylabel("score");
```



```
In [67]: 1 best_alpha = alphas[np.argmax(cv_scores)]
         2 best_alpha
```

Out[67]: 100.0

- It seems  $\alpha=100$  is the best choice here.

- General intuition: larger  $\alpha$  leads to smaller coefficients.
- Smaller coefficients mean the predictions are less sensitive to changes in the data.
- Hence less chance of overfitting (seeing big dependencies when you shouldn't).

## RidgeCV

BTW, because it's so common to want to tune  $\alpha$  with Ridge, sklearn provides a class called `RidgeCV`, which automatically tunes  $\alpha$  based on cross-validation.

```
In [68]: 1 ridgecv_pipe = make_pipeline(preprocessor, RidgeCV(alphas=alphas, cv=10)
         2 ridgecv_pipe.fit(X_train, y_train);
```

```
In [69]: 1 best_alpha = ridgecv_pipe.named_steps['ridgecv'].alpha_
         2 best_alpha
```

Out[69]: 100.0

## Let's examine the coefficients

```
In [70]: 1 lr_tuned = make_pipeline(preprocessor, Ridge(alpha=best_alpha))
         2 lr_tuned.fit(X_train, y_train)
         3 lr_preds = lr_tuned.predict(X_test)
         4 lr_preds[:10]
```

Out[70]: array([228728.1963872 , 104718.39905565, 155778.96723311, 246316.7111903  
1,  
127633.10676873, 243207.19441128, 304930.24461291, 145374.5943529  
5,  
157059.38983893, 128487.51979632])

```
In [71]: 1 lr_preds.max(), lr_preds.min()
```

Out[71]: (390726.10647423274, 30791.092505420726)

Let's get the feature names of the transformed data to interpret coefficients.

```
In [72]: 1 ohe_columns = list(
2         preprocessor.named_transformers_["pipeline-4"]
3         .named_steps["onehotencoder"]
4         .get_feature_names(categorical_features)
5     )
6 new_columns = numeric_features + ordinal_features_reg + ordinal_features
7
```

```
In [73]: 1 df = pd.DataFrame(
2         data={
3             "features": new_columns,
4             "coefficients": lr_tuned.named_steps["ridge"].coef_,
5         }
6     )
```

```
In [74]: 1 df.sort_values("coefficients", ascending=False)
```

```
Out[74]:
```

	features	coefficients
4	OverallQual	14484.902165
16	GrLivArea	11704.053037
70	Neighborhood_NridgHt	9662.969631
69	Neighborhood_NoRidge	9497.598615
36	BsmtQual	8073.088562
...	...	...
249	RoofMatl_ClyTile	-3992.399179
245	LandContour_Bnk	-5001.996997
62	Neighborhood_Gilbert	-5197.585536
59	Neighborhood_CollgCr	-5467.463086
61	Neighborhood_Edwards	-5796.508529

263 rows × 2 columns

So according to this model:

- As OverallQual feature gets bigger the housing price will get bigger.
- Presence of Neighborhood\_Edwards will result in smaller median house value.

```
In [75]: 1 x_train_enc[ 'Neighborhood_Edwards' ]
```

```
Out[75]: 302      0.0
          767      0.0
          429      0.0
          1139     0.0
          558      0.0
          ...
          1041     0.0
          1122     1.0
          1346     0.0
          1406     0.0
          1389     0.0
          Name: Neighborhood_Edwards, Length: 1314, dtype: float64
```

## Regression score functions

- We aren't doing classification anymore, so we can't just check for equality:

```
In [76]: 1 # This doesn't make sense:
          2 lr_tuned.predict(X_train) == y_train
```

```
Out[76]: 302      False
          767      False
          429      False
          1139      False
          558      False
          ...
          1041      False
          1122      False
          1346      False
          1406      False
          1389      False
          Name: SalePrice, Length: 1314, dtype: bool
```

```
In [77]: 1 y_train.values
```

```
Out[77]: array([205000, 160000, 175000, ..., 262500, 133000, 131000])
```

```
In [78]: 1 lr_tuned.predict(X_train)
```

```
Out[78]: array([212894.62756285, 178502.78223444, 189937.18327372, ...,
                245233.6751565 , 129863.13373552, 135439.89186716])
```

Processing math: 100% We need a score that reflects how right/wrong each prediction is.

There are a number of popular scoring functions for regression. We are going to look at some common metrics:

- mean squared error (MSE)
- $R^2$
- root mean squared error (RMSE)
- MAPE

See [sklearn documentation \(https://scikit-learn.org/stable/modules/model\\_evaluation.html#regression-metrics\)](https://scikit-learn.org/stable/modules/model_evaluation.html#regression-metrics) for more details.

## Mean squared error (MSE)

- A common metric is mean squared error:

```
In [79]: 1 preds = lr_tuned.predict(X_train)
```

```
In [80]: 1 np.mean((y_train - preds) ** 2)
```

```
Out[80]: 873230473.3636098
```

Perfect predictions would have MSE=0.

```
In [81]: 1 np.mean((y_train - y_train) ** 2)
```

```
Out[81]: 0.0
```

This is also implemented in sklearn:

```
In [82]: 1 from sklearn.metrics import mean_squared_error
2
3 mean_squared_error(y_train, preds)
```

```
Out[82]: 873230473.3636098
```

- MSE looks huge and unreasonable. There is an error of ~\$1 Billion!
- Is this score good or bad?

- Unlike classification, with regression **our target has units**.
- The target is in dollars, the mean squared error is in  $\text{\$dollars}^2$
- The score also depends on the scale of the targets.
- If we were working in cents instead of dollars, our MSE would be 10,000 times  $\text{\$(100}^2\text{\$)}$  higher!

```
In [83]: 1 np.mean((y_train * 100 - preds * 100) ** 2)
```

```
Out[83]: 8732304733636.098
```

## Root mean squared error or RMSE


- The MSE above is in  $\text{dollars}^2$ .
- A more relatable metric would be the root mean squared error, or RMSE

```
In [84]: 1 np.sqrt(mean_squared_error(y_train, lr_tuned.predict(X_train)))
```

```
Out[84]: 29550.473318774606
```

- Error of \$30,000 makes more sense.
- Can we dig deeper?

```
In [85]: 1 plt.scatter(y_train, lr_tuned.predict(X_train), alpha=0.3)
2 grid = np.linspace(y_train.min(), y_train.max(), 1000)
3 plt.plot(grid, grid, "--k")
4 plt.xlabel("true price")
5 plt.ylabel("predicted price");
```



- Here we can see a few cases where our prediction is way off.
- Is there something weird about those houses, perhaps? Outliers?
- Under the line means we're under-prediction, over the line means we're over-predicting.

Processing math: 100%



## $R^2$ (not in detail)

A common score is the  $R^2$

- This is the score that `sklearn` uses by default when you call `score()`:
- You can [read about it \(https://en.wikipedia.org/wiki/Coefficient\\_of\\_determination\)](https://en.wikipedia.org/wiki/Coefficient_of_determination) if interested.
- Intuition: similar to mean squared error, but flipped (higher is better), and normalized so the max is 1.

$$R^2(y, \hat{y}) = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

Key points:

- The maximum is 1 for perfect predictions
- Negative values are very bad: "worse than DummyRegressor" (very bad)

(optional) Warning: MSE is "reversible" but  $R^2$  is not:

```
In [86]: 1 mean_squared_error(y_train, preds)
```

```
Out[86]: 873230473.3636098
```

```
In [87]: 1 mean_squared_error(preds, y_train)
```

```
Out[87]: 873230473.3636098
```

```
In [88]: 1 r2_score(y_train, preds)
```

```
Out[88]: 0.8601212294857903
```

```
In [89]: 1 r2_score(preds, y_train)
```

```
Out[89]: 0.827962225882707
```

- When you call `fit` it minimizes MSE / maximizes  $R^2$  (or something like that) by default.
- Just like in classification, this isn't always what you want!!

## MAPE

- We got an RMSE of ~\$30,000 before.

Question: Is an error of \$30,000 acceptable?


```
In [90]: 1 np.sqrt(mean_squared_error(y_train, lr_tuned.predict(X_train)))
```

```
Out[90]: 29550.473318774606
```

- For a house worth \$600k, it seems reasonable! That's 5% error.
- For a house worth \$60k, that is terrible. It's 50% error.

We have both of these cases in our dataset.

```
In [91]: 1 plt.hist(y_train, bins=100);
```



How about looking at percent error?

```
In [92]: 1 pred_train = lr_tuned.predict(X_train)
          2 percent_errors = (pred_train - y_train) / y_train * 100.0
          3 percent_errors
```

```
Out[92]: 302      3.851038
          767      11.564239
          429      8.535533
          1139    -16.371069
          558      17.177968
          ...
          1041    -0.496571
          1122   -28.696351
          1346    -6.577648
          1406    -2.358546
          1389      3.389230
          Name: SalePrice, Length: 1314, dtype: float64
```


These are both positive (predict too high) and negative (predict too low).

We can look at the absolute percent error:

```
In [93]: 1 np.abs(percent_errors)
```

```
Out[93]: 302      3.851038
          767      11.564239
          429      8.535533
          1139     16.371069
          558      17.177968
          ...
          1041      0.496571
          1122     28.696351
          1346      6.577648
          1406      2.358546
          1389      3.389230
          Name: SalePrice, Length: 1314, dtype: float64
```

```
In [94]: 1 plt.hist(np.abs(percent_errors), bins=100);
```



And, like MSE, we can take the average over examples. This is called mean absolute percent error (MAPE).

```
In [95]: 1 def mape(true, pred):  
2     return 100.0 * np.mean(np.abs((pred - true) / true))
```

```
In [96]: 1 mape(y_train, pred_train)
```

```
Out[96]: 10.093121294225265
```

- Ok, this is quite interpretable.
- On average, we have around 10% error.

## Transforming the targets

- Does `.fit()` know we care about MAPE?
- No, it doesn't. Why are we minimizing MSE (or something similar) if we care about MAPE??
- When minimizing MSE, the expensive houses will dominate because they have the biggest error.
- Which is better for RMSE?


- Example 1: Truth: \$50k, Prediction: \ \$100k
- Example 2: Truth: \$500k, Prediction: \ \$550k
- RMSE: \$50k
- MAPE: 45%

#### Model B


- Example 1: Truth: \$50k, Prediction: \ \$60k
- Example 2: Truth: \$500k, Prediction: \ \$600k
- RMSE: \$71k
- MAPE: 20%

- How can we get `.fit()` to think about MAPE?
- A common practice which tends to work is log transforming the targets.
- That is, transform  $y \rightarrow \log(y)$ .

```
In [97]: 1 plt.hist(y_train, bins=100);
```



```
In [98]: 1 plt.hist(np.log10(y_train), bins=100);
```



We can incorporate this in our pipeline using `sklearn`.

```
In [99]: 1 from sklearn.compose import TransformedTargetRegressor
```

```
In [100]: 1 ttr = TransformedTargetRegressor(
2     Ridge(alpha=best_alpha), func=np.log1p, inverse_func=np.expml
3 ) # transformer for log transforming the target
4 ttr_pipe = make_pipeline(preprocessor, ttr)
```

```
In [101]: 1 ttr_pipe.fit(X_train, y_train); # y_train automatically transformed
```

```
In [102]: 1 ttr_pipe.predict(X_train) # predictions automatically un-transformed
```

```
Out[102]: array([221355.29528077, 170663.43286226, 182608.09768702, ...,
248575.94877669, 132148.9047652 , 133262.17638244])
```

```
In [103]: 1 mape(y_test, ttr_pipe.predict(X_test))
```

```
Out[103]: 7.808600924240852
```

We reduced MAPE from ~10% to ~8% with this trick!

## Different scoring functions with `cross_validate`

- Let's try using MSE instead of the default  $R^2$  score.

```
In [104]: 1 pd.DataFrame(
2           cross_validate(
3             lr_tuned,
4             X_train,
5             y_train,
6             return_train_score=True,
7             scoring="neg_mean_squared_error",
8           )
9         )
```

```
Out[104]:
```

	fit_time	score_time	test_score	train_score
0	0.070120	0.027592	-7.060346e+08	-9.383069e+08
1	0.073525	0.027878	-1.239851e+09	-8.267971e+08
2	0.071608	0.048744	-1.125125e+09	-8.763019e+08
3	0.111845	0.030674	-9.819320e+08	-8.847908e+08
4	0.072874	0.026326	-2.268434e+09	-7.397199e+08

```
In [105]: 1 def mape(true, pred):
2           return 100.0 * np.mean(np.abs((pred - true) / true))
3
4
5 # make a scorer function that we can pass into cross-validation
6 mape_scorer = make_scorer(mape, greater_is_better=True)
7
8 pd.DataFrame(
9     cross_validate(
10         lr_tuned, X_train, y_train, return_train_score=True, scoring=ma
11     )
12 )
```

```
Out[105]:
```

	fit_time	score_time	test_score	train_score
0	0.077079	0.025021	9.699277	10.407124
1	0.076896	0.025295	10.803043	9.966190
2	0.094964	0.030100	11.836195	10.180734
3	0.090991	0.037569	10.784686	10.247198
4	0.081499	0.029051	12.196718	9.828607

```
In [106]: 1 scoring = {
2         "r2": "r2",
3         "mape_scorer": mape_scorer,
4         "neg_root_mean_square_error": "neg_root_mean_squared_error",
5         "neg_mean_squared_error": "neg_mean_squared_error",
6     }
7
8     pd.DataFrame(
9         cross_validate(lr_tuned, X_train, y_train, return_train_score=True,
10                        ).T
```

```
Out[106]:
```

	0	1	2	3	
fit_time	8.490086e-02	8.591008e-02	8.298492e-02	8.994293e-02	8.87
score_time	3.135419e-02	3.481293e-02	3.112602e-02	3.396201e-02	3.49
test_r2	8.668969e-01	8.200460e-01	8.262644e-01	8.511854e-01	6.10
train_r2	8.551369e-01	8.636241e-01	8.579735e-01	8.561893e-01	8.83
test_mape_scorer	9.699277e+00	1.080304e+01	1.183620e+01	1.078469e+01	1.215
train_mape_scorer	1.040712e+01	9.966190e+00	1.018073e+01	1.024720e+01	9.821
test_neg_root_mean_square_error	-2.657131e+04	-3.521152e+04	-3.354288e+04	-3.133579e+04	-4.762
train_neg_root_mean_square_error	-3.063179e+04	-2.875408e+04	-2.960240e+04	-2.974543e+04	-2.715
test_neg_mean_squared_error	-7.060346e+08	-1.239851e+09	-1.125125e+09	-9.819320e+08	-2.261
train_neg_mean_squared_error	-9.383069e+08	-8.267971e+08	-8.763019e+08	-8.847908e+08	-7.397

```
In [107]: 1 mape(y_test, lr_tuned.predict(X_test))
```

```
Out[107]: 9.496387589496008
```

## Using regression metrics with `scikit-learn`

- In `sklearn` you will notice that it has negative version of the metrics above (e.g., `neg_mean_squared_error`, `neg_root_mean_squared_error`).
- The reason for this is that scores return a value to maximize, the higher the better.
- If you define your own scorer function and if you do not want this interpretation, you can set the `greater_is_better` parameter to `False`

## Questions for class discussion

### True/False

1. Price per square foot would be a good feature to add in our `x`.
2. The `alpha` hyperparameter of `Ridge` has similar interpretation of `c` hyperparameter of `LogisticRegression`; higher `alpha` means more complex model.

Processing math: 100%



3. In regression, one should use MAPE instead of MSE when relative (percent) error matters more than absolute error.
4. A lower RMSE value indicates a better model.
5. We can still use precision and recall for regression problems but now we have other

## Summary

- House prices dataset target is price, which is numeric -> regression rather than classification
- There are corresponding versions of all the tools we used:
  - `DummyClassifier` -> `DummyRegressor`
  - `LogisticRegression` -> `Ridge`
- Ridge hyperparameter `alpha` is like `LogisticRegression` hyperparameter `C`, but opposite meaning
- We'll avoid `LinearRegression` in this course.

- Scoring metrics
- $R^2$  is the default `.score()`, it is unitless, 0 is bad, 1 is best
- MSE (mean squared error) is in units of target squared, hard to interpret; 0 is best
- RMSE (root mean squared error) is in the same units as the target; 0 is best
- MAPE (mean average percent error) is unitless; 0 is best, 100 is bad

In [ ]:

1