

CPSC 330

Applied Machine Learning

Lecture 20: Survival analysis

UBC 2022-23

Instructor: Mathias Lécuyer

Imports

```
In [1]: 1 import matplotlib.pyplot as plt
2 import numpy as np
3 import pandas as pd
4 from sklearn.compose import ColumnTransformer, make_column_transformer
5 from sklearn.dummy import DummyClassifier
6 from sklearn.ensemble import RandomForestClassifier, RandomForestRegressor
7 from sklearn.impute import SimpleImputer
8 from sklearn.linear_model import LogisticRegression, Ridge
9 from sklearn.metrics import confusion_matrix, plot_confusion_matrix
10 from sklearn.model_selection import (
11     cross_val_predict,
12     cross_val_score,
13     cross_validate,
14     train_test_split,
15 )
16 from sklearn.pipeline import Pipeline, make_pipeline
17 from sklearn.preprocessing import (
18     FunctionTransformer,
19     OneHotEncoder,
20     OrdinalEncoder,
21     StandardScaler,
22 )
23
24 plt.rcParams["font.size"] = 16
25
26 # does lifelines try to mess with this?
27 pd.options.display.max_rows = 10
```

```
In [2]: 1 import lifelines
```

Learning objectives

- Explain the problem with treating right-censored data the same as "regular" data.

- Determine whether survival analysis is an appropriate tool for a given problem.
- Apply survival analysis in Python using the `lifelines` package.
- Interpret a survival curve, such as the Kaplan-Meier curve.
- Interpret the coefficients of a fitted Cox proportional hazards model.
- Make predictions for existing individuals and interpret these predictions.

Customer churn: our standard approach

- In hw5 you looked at a dataset about [customer churn](https://en.wikipedia.org/wiki/Customer_attrition) (https://en.wikipedia.org/wiki/Customer_attrition).
- In hw5, the dataset was interesting because it's unbalanced (most customers stay). We used typical binary classification approach on the dataset.
- Today we'll look at a different customer churn [dataset](https://www.kaggle.com/blastchar/telco-customer-churn) (<https://www.kaggle.com/blastchar/telco-customer-churn>), because it has a feature we need - time!
- We'll explore the time aspect of the dataset today.

```
In [5]: 1 df = pd.read_csv("../data/WA_Fn-UseC_-Telco-Customer-Churn.csv")
        2 train_df, test_df = train_test_split(df, random_state=123)
        3 train_df.head()
```

```
Out[5]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines
6464	4726-DLWQN	Male	1	No	No	50	Yes	Yes
5707	4537-DKTAL	Female	0	No	No	2	Yes	No
3442	0468-YRPXN	Male	0	No	No	29	Yes	No
3932	1304-NECVQ	Female	1	No	No	2	Yes	Yes
6124	7153-CHRBV	Female	0	Yes	Yes	57	Yes	No

5 rows × 21 columns

We can treat this as a binary classification problem where we want to predict `Churn` (yes/no) from these other columns.

```
In [6]: 1 train_df.shape
```

```
Out[6]: (5282, 21)
```

```
In [7]: 1 train_df["Churn"].value_counts()
```

```
Out[7]: No      3912
        Yes     1370
        Name: Churn, dtype: int64
```

In [8]: 1 train_df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 5282 entries, 6464 to 3582
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   customerID            5282 non-null   object
 1   gender                5282 non-null   object
 2   SeniorCitizen         5282 non-null   int64
 3   Partner               5282 non-null   object
 4   Dependents            5282 non-null   object
 5   tenure                5282 non-null   int64
 6   PhoneService          5282 non-null   object
 7   MultipleLines          5282 non-null   object
 8   InternetService       5282 non-null   object
 9   OnlineSecurity        5282 non-null   object
10   OnlineBackup          5282 non-null   object
11   DeviceProtection      5282 non-null   object
12   TechSupport           5282 non-null   object
13   StreamingTV           5282 non-null   object
14   StreamingMovies       5282 non-null   object
15   Contract              5282 non-null   object
16   PaperlessBilling      5282 non-null   object
17   PaymentMethod         5282 non-null   object
18   MonthlyCharges        5282 non-null   float64
19   TotalCharges          5282 non-null   object
20   Churn                 5282 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 907.8+ KB
```

Question: Does this mean there is no missing data?

Ok, let's try our usual approach:

In [9]: 1 train_df["SeniorCitizen"].value_counts()

```
Out[9]: 0    4430
        1     852
        Name: SeniorCitizen, dtype: int64
```

```
In [10]: 1 numeric_features = ["tenure", "MonthlyCharges", "TotalCharges"]
2 drop_features = ["customerID"]
3 passthrough_features = ["SeniorCitizen"]
4 target_column = ["Churn"]
5 # the rest are categorical
6 categorical_features = list(
7     set(train_df.columns)
8     - set(numeric_features)
9     - set(passthrough_features)
10    - set(drop_features)
11    - set(target_column)
12 )
```

```
In [11]: 1 preprocessor = make_column_transformer(
2     (StandardScaler(), numeric_features),
3     (OneHotEncoder(), categorical_features),
4     ("passthrough", passthrough_features),
5     ("drop", drop_features),
6 )
```

```
In [12]: 1 preprocessor.fit(train_df);
```

```
-----
--
ValueError                                Traceback (most recent call last)
Cell In[12], line 1
----> 1 preprocessor.fit(train_df);

File /opt/anaconda3/envs/cpsc330/lib/python3.10/site-packages/sklearn/com
pose/_column_transformer.py:657, in ColumnTransformer.fit(self, X, y)
    639 """Fit all transformers using X.
    640
    641 Parameters
    (...)
    653     This estimator.
    654 """
    655 # we use fit_transform to make sure to set sparse_output_ (for wh
ich we
    656 # need the transformed data) to have consistent output type in pr
edict
    657 self.fit_transform(X, y)
```

Hmmm, one of the numeric features is causing problems?

In [13]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                 7043 non-null   object
2   SeniorCitizen         7043 non-null   int64
3   Partner               7043 non-null   object
4   Dependents            7043 non-null   object
5   tenure                7043 non-null   int64
6   PhoneService          7043 non-null   object
7   MultipleLines         7043 non-null   object
8   InternetService       7043 non-null   object
9   OnlineSecurity        7043 non-null   object
10  OnlineBackup          7043 non-null   object
11  DeviceProtection      7043 non-null   object
12  TechSupport           7043 non-null   object
13  StreamingTV           7043 non-null   object
14  StreamingMovies       7043 non-null   object
15  Contract              7043 non-null   object
16  PaperlessBilling      7043 non-null   object
17  PaymentMethod         7043 non-null   object
18  MonthlyCharges        7043 non-null   float64
19  TotalCharges          7043 non-null   object
20  Churn                 7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

Oh, looks like `TotalCharges` is not a numeric type. What if we change the type of this column to float?

```
In [14]: 1 train_df["TotalCharges"] = train_df["TotalCharges"].astype(float)
```

```

-----
--
ValueError                                Traceback (most recent call las
t)
Cell In[14], line 1
----> 1 train_df["TotalCharges"] = train_df["TotalCharges"].astype(float)

File /opt/anaconda3/envs/cpsc330/lib/python3.10/site-packages/pandas/core/generic.py:6240, in NDFrame.astype(self, dtype, copy, errors)
    6233         results = [
    6234             self.iloc[:, i].astype(dtype, copy=copy)
    6235             for i in range(len(self.columns))
    6236         ]
    6237     else:
    6238         # else, only a single dtype is given
-> 6240         new_data = self._mgr.astype(dtype=dtype, copy=copy, errors=er
rors)
    6241         return self._constructor(new_data).__finalize__(self, method
="astype")
    6242 # GH 33113: handle empty frame or series

File /opt/anaconda3/envs/cpsc330/lib/python3.10/site-packages/pandas/core/internals/managers.py:448, in BaseBlockManager.astype(self, dtype, copy, errors)
    447 def astype(self: T, dtype, copy: bool = False, errors: str = "rai
se") -> T:
-> 448         return self.apply("astype", dtype=dtype, copy=copy, errors=er
rors)

File /opt/anaconda3/envs/cpsc330/lib/python3.10/site-packages/pandas/core/internals/managers.py:352, in BaseBlockManager.apply(self, f, align_key
s, ignore_failures, **kwargs)
    350         applied = b.apply(f, **kwargs)
    351     else:
-> 352         applied = getattr(b, f)(**kwargs)
    353 except (TypeError, NotImplementedError):
    354     if not ignore_failures:

File /opt/anaconda3/envs/cpsc330/lib/python3.10/site-packages/pandas/core/internals/blocks.py:526, in Block.astype(self, dtype, copy, errors)
    508 """
    509 Coerce to the new dtype.
    510
    511 (...)
    522 Block
    523 """
    524 values = self.values
-> 526 new_values = astype_array_safe(values, dtype, copy=copy, errors=e
rrors)
    527 new_values = maybe_coerce_values(new_values)
    528 newb = self.make_block(new_values)

File /opt/anaconda3/envs/cpsc330/lib/python3.10/site-packages/pandas/core/dtypes/astype.py:299, in astype_array_safe(values, dtype, copy, errors)
    296         return values.copy()
    297     try:
-> 299         new_values = astype_array(values, dtype, copy=copy)

```

```

300 except (ValueError, TypeError):
301     # e.g. astype_nansafe can fail on object-dtype of strings
302     # trying to convert to float
303     if errors == "ignore":

```

File /opt/anaconda3/envs/cpsc330/lib/python3.10/site-packages/pandas/core/dtypes/astype.py:230, in astype_array(values, dtype, copy)

```

227     values = values.astype(dtype, copy=copy)
229 else:
--> 230     values = astype_nansafe(values, dtype, copy=copy)
232 # in pandas we don't store numpy str dtypes, so convert to object
233 if isinstance(dtype, np.dtype) and issubclass(values.dtype.type,
str):

```

File /opt/anaconda3/envs/cpsc330/lib/python3.10/site-packages/pandas/core/dtypes/astype.py:170, in astype_nansafe(arr, dtype, copy, skipna)

```

166     raise ValueError(msg)
168 if copy or is_object_dtype(arr.dtype) or is_object_dtype(dtype):
169     # Explicit copy, or required since NumPy can't view from / to
object.
--> 170     return arr.astype(dtype, copy=True)
172 return arr.astype(dtype, copy=copy)

```

ValueError: could not convert string to float: ' '

Argh!!

```

In [15]: 1 for val in train_df["TotalCharges"]:
2         try:
3             float(val)
4         except ValueError:
5             print(val)

```

Any ideas?

Well, it turns out we can't see those problematic values because they are whitespace!


```
In [16]: 1 for val in train_df["TotalCharges"]:
          2     try:
          3         float(val)
          4     except ValueError:
          5         print('%s' % val)
```

```
" "
" "
" "
" "
" "
" "
" "
```

Let's replace the whitespaces with NaNs.

```
In [17]: 1 train_df = train_df.assign(
          2     TotalCharges=train_df["TotalCharges"].replace(" ", np.nan).astype(f
          3 )
          4 test_df = test_df.assign(
          5     TotalCharges=test_df["TotalCharges"].replace(" ", np.nan).astype(fl
          6 )
```

```
In [18]: 1 train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 5282 entries, 6464 to 3582
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            5282 non-null   object
1   gender                5282 non-null   object
2   SeniorCitizen         5282 non-null   int64
3   Partner               5282 non-null   object
4   Dependents            5282 non-null   object
5   tenure               5282 non-null   int64
6   PhoneService          5282 non-null   object
7   MultipleLines         5282 non-null   object
8   InternetService       5282 non-null   object
9   OnlineSecurity        5282 non-null   object
10  OnlineBackup          5282 non-null   object
11  DeviceProtection      5282 non-null   object
12  TechSupport           5282 non-null   object
13  StreamingTV           5282 non-null   object
14  StreamingMovies       5282 non-null   object
15  Contract              5282 non-null   object
16  PaperlessBilling      5282 non-null   object
17  PaymentMethod         5282 non-null   object
18  MonthlyCharges        5282 non-null   float64
19  TotalCharges          5274 non-null   float64
20  Churn                 5282 non-null   object
dtypes: float64(2), int64(2), object(17)
memory usage: 907.8+ KB
```

But now we are going to have missing values and we need to include imputation for numeric features in our preprocessor.

```
In [19]: 1 preprocessor = make_column_transformer(
2         (
3             make_pipeline(SimpleImputer(strategy="median"), StandardScaler(
4                 numeric_features,
5             ),
6             (OneHotEncoder(handle_unknown="ignore"), categorical_features),
7             ("passthrough", passthrough_features),
8             ("drop", drop_features),
9         )
```

Now let's try that again...

```
In [20]: 1 preprocessor.fit(train_df);
```

It worked! Let's get the column names of the transformed data from the column transformer.

```
In [22]: 1 new_columns = (
2         numeric_features
3         + preprocessor.named_transformers_["onehotencoder"]
4         .get_feature_names_out(categorical_features)
5         .tolist()
6         + passthrough_features
7     )
```

```
In [23]: 1 X_train_enc = pd.DataFrame(
2         preprocessor.transform(train_df), index=train_df.index, columns=new
3     )
4 X_test_enc = pd.DataFrame(
5     preprocessor.transform(train_df), index=train_df.index, columns=new
6     )
```

```
In [24]: 1 X_train_enc.head()
```

```
Out[24]:
```

	tenure	MonthlyCharges	TotalCharges	TechSupport_No	TechSupport_No internet service	TechSupport_Yes
6464	0.707712	0.185175	0.513678	1.0	0.0	0.0
5707	-1.248999	-0.641538	-0.979562	1.0	0.0	0.0
3442	-0.148349	1.133562	0.226789	0.0	0.0	1.0
3932	-1.248999	0.458524	-0.950696	1.0	0.0	0.0
6124	0.993065	-0.183179	0.433814	0.0	0.0	1.0

5 rows × 45 columns

Before we look into survival analysis, let's just treat it as a binary classification model where we want to predict whether a customer churned or not.

```
In [25]: 1 results = {}
```

```
In [26]: 1 def mean_std_cross_val_scores(model, X_train, y_train, **kwargs):
2         """
3         Returns mean and std of cross validation
4
5         Parameters
6         -----
7         model :
8             scikit-learn model
9         X_train : numpy array or pandas DataFrame
10            X in the training data
11         y_train :
12            y in the training data
13
14         Returns
15         -----
16            pandas Series with mean scores from cross_validation
17         """
18
19         scores = cross_validate(model, X_train, y_train, **kwargs)
20
21         mean_scores = pd.DataFrame(scores).mean()
22         std_scores = pd.DataFrame(scores).std()
23         out_col = []
24
25         for i in range(len(mean_scores)):
26             out_col.append((f"%0.3f (+/- %0.3f)" % (mean_scores[i], std_scores[i])))
27
28         return pd.Series(data=out_col, index=mean_scores.index)
```

```
In [27]: 1 X_train = train_df.drop(columns=["Churn"])
2         X_test = test_df.drop(columns=["Churn"])
3
4         y_train = train_df["Churn"]
5         y_test = test_df["Churn"]
```

DummyClassifier

```
In [28]: 1 dc = DummyClassifier()
```

```
In [29]: 1 results["dummy"] = mean_std_cross_val_scores(
2         dc, X_train, y_train, return_train_score=True
3     )
4     pd.DataFrame(results)
```

```
Out[29]:
```

	dummy
fit_time	0.003 (+/- 0.001)
score_time	0.001 (+/- 0.000)
test_score	0.741 (+/- 0.000)
train_score	0.741 (+/- 0.000)

LogisticRegression

```
In [30]: 1 lr = make_pipeline(preprocessor, LogisticRegression(max_iter=1000))
```

```
In [31]: 1 results["logistic regression"] = mean_std_cross_val_scores(
2         lr, X_train, y_train, return_train_score=True
3     )
4     pd.DataFrame(results)
```

```
Out[31]:
```

	dummy	logistic regression
fit_time	0.003 (+/- 0.001)	0.055 (+/- 0.008)
score_time	0.001 (+/- 0.000)	0.006 (+/- 0.000)
test_score	0.741 (+/- 0.000)	0.804 (+/- 0.013)
train_score	0.741 (+/- 0.000)	0.809 (+/- 0.002)

```
In [32]: 1 confusion_matrix(y_train, cross_val_predict(lr, X_train, y_train))
```

```
Out[32]: array([[3516,  396],
                [ 637,  733]])
```

RandomForestClassifier

```
In [33]: 1 rf = make_pipeline(preprocessor, RandomForestClassifier())
```

```
In [34]: 1 results["random forest"] = mean_std_cross_val_scores(
2         rf, X_train, y_train, return_train_score=True
3     )
4     pd.DataFrame(results)
```

```
Out[34]:
```

	dummy	logistic regression	random forest
fit_time	0.003 (+/- 0.001)	0.055 (+/- 0.008)	0.287 (+/- 0.014)
score_time	0.001 (+/- 0.000)	0.006 (+/- 0.000)	0.020 (+/- 0.001)
test_score	0.741 (+/- 0.000)	0.804 (+/- 0.013)	0.790 (+/- 0.012)
train_score	0.741 (+/- 0.000)	0.809 (+/- 0.002)	0.998 (+/- 0.000)

```
In [35]: 1 confusion_matrix(y_train, cross_val_predict(rf, X_train, y_train))
```

```
Out[35]: array([[3518,  394],
                [ 735,  635]])
```

- This is was we did in hw5.
- What's wrong with this approach?

And now the rest of the class is about what is wrong with what we just did!

Censoring and survival analysis

Time to event and censoring

Imagine that you want to analyze *the time until an event occurs*. For example,

- the time until a disease kills its host.
- the time until a piece of equipment breaks.
- the time that someone unemployed will take to land a new job.
- the time until a customer leaves a subscription service (this dataset).

In our example, instead of predicting the binary label churn or no churn, it will be more useful to estimate when the customer is likely to churn (the time until churn happens) so that we can take some action.

```
In [36]: 1 train_df[["tenure"]].head()
```

```
Out[36]:
```

	tenure
6464	50
5707	2
3442	29
3932	2
6124	57

The tenure column is the number of months the customer has stayed with the company.

Although this branch of statistics is usually referred to as **Survival Analysis**, the event in question does not need to be related to actual "survival". The important thing is to understand that we are interested in **the time until something happens**, or whether or not something will happen in a certain time frame.

Question: But why is this different? Can't you just use the techniques you learned so far (e.g., regression models) to predict the time? Take a minute to think about this.

The answer would be yes if you could observe the actual time in all occurrences, but you usually cannot. Frequently, there will be some kind of **censoring** which will not allow you to observe the exact time that the event happened for all units/individuals that are being studied.

```
In [37]: 1 train_df[["tenure", "Churn"]].head()
```

```
Out[37]:
```

	tenure	Churn
6464	50	No
5707	2	No
3442	29	No
3932	2	Yes
6124	57	No

- What this means is that we **don't have correct target values** to train or test our model.
- This is a problem!

Let's consider some approaches to deal with this censoring issue.

Approach 1: Only consider the examples where "Churn"=Yes

Let's just consider the cases *for which we have the time*, to obtain the average subscription length.

```
In [38]: 1 train_df_churn = train_df.query(
2         "Churn == 'Yes'"
3     ) # Consider only examples where the customers churned.
4 test_df_churn = test_df.query(
5         "Churn == 'Yes'"
6     ) # Consider only examples where the customers churned.
7 train_df_churn.head()
```

```
Out[38]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines
3932	1304-NECVQ	Female	1	No	No	2	Yes	Yes
301	8098-LLAZX	Female	1	No	No	4	Yes	Yes
5540	3803-KMQFW	Female	0	Yes	Yes	1	Yes	No
4084	2777-PHDEI	Female	0	No	No	1	Yes	No
3272	6772-KSATR	Male	0	No	No	1	Yes	Yes

5 rows × 21 columns

```
In [39]: 1 train_df.shape
```

```
Out[39]: (5282, 21)
```

```
In [40]: 1 train_df_churn.shape
```

```
Out[40]: (1370, 21)
```

```
In [41]: 1 numeric_features
```

```
Out[41]: ['tenure', 'MonthlyCharges', 'TotalCharges']
```

```
In [42]: 1 preprocessing_notenure = make_column_transformer(
2         (
3             make_pipeline(SimpleImputer(strategy="median"), StandardScaler(
4                 numeric_features[1:], # Getting rid of the tenure column
5             ),
6             (OneHotEncoder(handle_unknown="ignore"), categorical_features),
7             ("passthrough", passthrough_features),
8         )
```

```
In [43]: 1 tenure_lm = make_pipeline(preprocessing_notenure, Ridge())
2
3 tenure_lm.fit(train_df_churn.drop(columns=["tenure"]), train_df_churn["
```

```
In [44]: 1 pd.DataFrame(  
2         tenure_lm.predict(test_df_churn.drop(columns=["tenure"]))[:10],  
3         columns=["tenure_predictions"],  
4     )
```

```
Out[44]:
```

	tenure_predictions
0	5.062449
1	13.198645
2	11.859455
3	5.865562
4	58.154842
5	3.757932
6	18.932070
7	7.720893
8	36.818041
9	7.263541

What will be wrong with our estimated survival times? Will they be too low or too high?

On average they will be **underestimates** (too small), because we are ignoring the currently subscribed (un-churned) customers. Our dataset is a biased sample of those who churned within the time window of the data collection. Long-time subscribers were more likely to be removed from the dataset! This is a common mistake - see the [Calling Bullshit video](https://www.youtube.com/watch?v=ITWQ5psx9Sw) (<https://www.youtube.com/watch?v=ITWQ5psx9Sw>) I posted on the README!

Approach 2: Assume everyone churns right now

Assume everyone churns right now - in other words, use the original dataset.


```
In [45]: 1 train_df[["tenure", "Churn"]].head()
```

```
Out[45]:
```

	tenure	Churn
6464	50	No
5707	2	No
3442	29	No
3932	2	Yes
6124	57	No

```
In [46]: 1 tenure_lm.fit(train_df.drop(columns=["tenure"]), train_df["tenure"]);
```

```
In [47]: 1 pd.DataFrame(  
2     tenure_lm.predict(test_df_churn.drop(columns=["tenure"]))[:10],  
3     columns=["tenure_predictions"],  
4 )
```

```
Out[47]:
```

	tenure_predictions
0	6.400047
1	20.220392
2	22.332746
3	12.825470
4	59.885968
5	7.075453
6	17.731498
7	10.407862
8	38.425365
9	10.854500

What will be wrong with our estimated survival time?

```
In [48]: 1 train_df[["tenure", "Churn"]].head()
```

```
Out[48]:
```

	tenure	Churn
6464	50	No
5707	2	No
3442	29	No
3932	2	Yes
6124	57	No

It will be an **underestimate** again. For those still subscribed, while we did not remove them, we recorded a total tenure shorter than in reality, because they will keep going for some amount of time.

Approach 3: Survival analysis

Deal with this properly using [survival analysis \(https://en.wikipedia.org/wiki/Survival_analysis\)](https://en.wikipedia.org/wiki/Survival_analysis).

- You may learn about this in a statistics course.
- We will use the `lifelines` package in Python and will not go into the math/stats of how it works.

```
In [49]: 1 train_df[["tenure", "Churn"]].head()
```


```
Out[49]:
```

	tenure	Churn
6464	50	No
5707	2	No
3442	29	No
3932	2	Yes
6124	57	No

Types of questions we might want to answer:

1. How long do customers stay with the service?
2. For a particular customer, can we predict how long they might stay with the service?
3. What factors influence a customer's churn time?

Break (5 min)



Kaplan-Meier survival curve

Before we do anything further, I want to modify our dataset slightly:

1. I'm going to drop the `TotalCharges` (yes, after all that work fixing it) because it's a bit of a strange feature.
 - Its value actually changes over time, but we only have the value at the end.
 - We still have `MonthlyCharges` .
2. I'm going to not scale the `tenure` column, since it will be convenient to keep it in its original units of months.

Just for our sanity, I'm redefining the features.

```
In [50]: 1 numeric_features = ["MonthlyCharges"]
2 drop_features = ["customerID", "TotalCharges"]
3 passthrough_features = ["tenure", "SeniorCitizen"] # don't want to sca
4 target_column = ["Churn"]
5 # the rest are categorical
6 categorical_features = list(
7     set(train_df.columns)
8     - set(numeric_features)
9     - set(passthrough_features)
10    - set(drop_features)
11    - set(target_column)
12 )
```

```
In [51]: 1 preprocessing_final = make_column_transformer(
2     (
3         FunctionTransformer(lambda x: x == "Yes"),
4         target_column,
5     ), # because we need it in this format for lifelines package
6     ("passthrough", passthrough_features),
7     (StandardScaler(), numeric_features),
8     (OneHotEncoder(handle_unknown="ignore", sparse=False), categorical_
9     ("drop", drop_features),
10 )
```

```
In [52]: 1 preprocessing_final.fit(train_df);
```

Let's get the column names of the columns created by our column transformer.

```
In [54]: 1 new_columns = (
2     target_column
3     + passthrough_features
4     + numeric_features
5     + preprocessing_final.named_transformers_["onehotencoder"]
6     .get_feature_names_out(categorical_features)
7     .tolist()
8 )
```

```
In [55]: 1 train_df_surv = pd.DataFrame(
2     preprocessing_final.transform(train_df), index=train_df.index, colu
3 )
4 test_df_surv = pd.DataFrame(
5     preprocessing_final.transform(test_df), index=test_df.index, column
6 )
```

```
In [56]: 1 train_df_surv.head()
```

Out[56]:

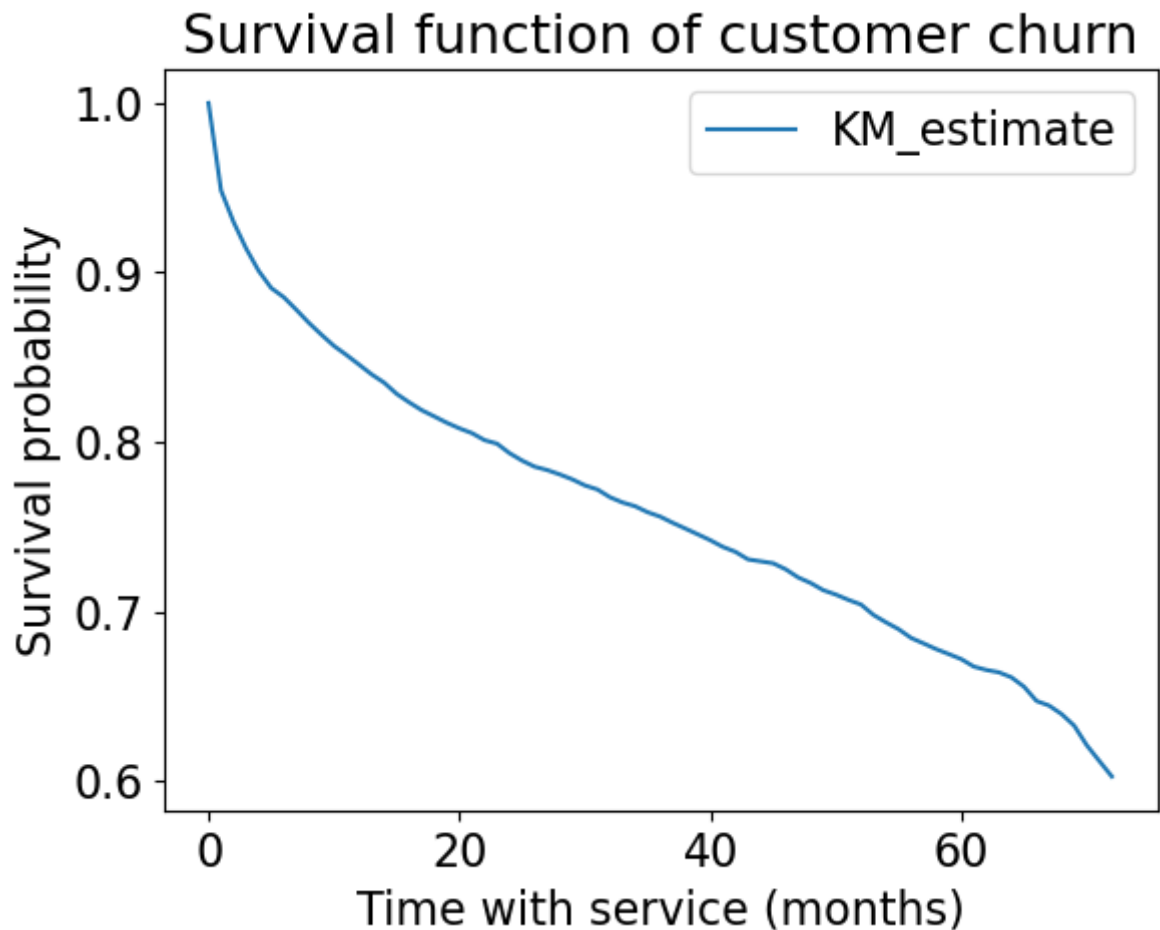
	Churn	tenure	SeniorCitizen	MonthlyCharges	TechSupport_No	TechSupport_No internet service	TechSupport
6464	0.0	50.0	1.0	0.185175	1.0	0.0	
5707	0.0	2.0	0.0	-0.641538	1.0	0.0	
3442	0.0	29.0	0.0	1.133562	0.0	0.0	
3932	1.0	2.0	1.0	0.458524	1.0	0.0	
6124	0.0	57.0	0.0	-0.183179	0.0	0.0	

5 rows × 45 columns

- We'll start with a model called `KaplanMeierFitter` from `lifelines` package to get a Kaplan Meier curve.
- For this model we only use two columns: `tenure` and `churn`.
- We do not use any other features.

```
In [57]: 1 kmf = lifelines.KaplanMeierFitter()
          2 kmf.fit(train_df_surv["tenure"], train_df_surv["Churn"]);
```

```
In [58]: 1 kmf.survival_function_.plot()  
2 plt.title("Survival function of customer churn")  
3 plt.xlabel("Time with service (months)")  
4 plt.ylabel("Survival probability");
```



- What is this plot telling us?
- It shows the probability of survival over time.
- For example, after 20 months the probability of survival is ~0.8.
- Over time it's going down.

```
In [64]: 1 np.mean(y_train == "No")
```

Out[64]: 0.7406285497917455

What's the average tenure?

```
In [65]: 1 np.mean(train_df_surv["tenure"])
```

Out[65]: 32.6391518364256

What's the average tenure of the people who churned?

```
In [66]: 1 np.mean(train_df_surv.query("Churn == 1.0")["tenure"])
```

```
Out[66]: 17.854744525547446
```


What's the average tenure of the people who did not churn?

```
In [67]: 1 np.mean(train_df_surv.query("Churn == 0.0")["tenure"])
```

```
Out[67]: 37.816717791411044
```

- Let's look at the histogram of number of people who have not churned.
- The key point here is that people *joined at different times*.

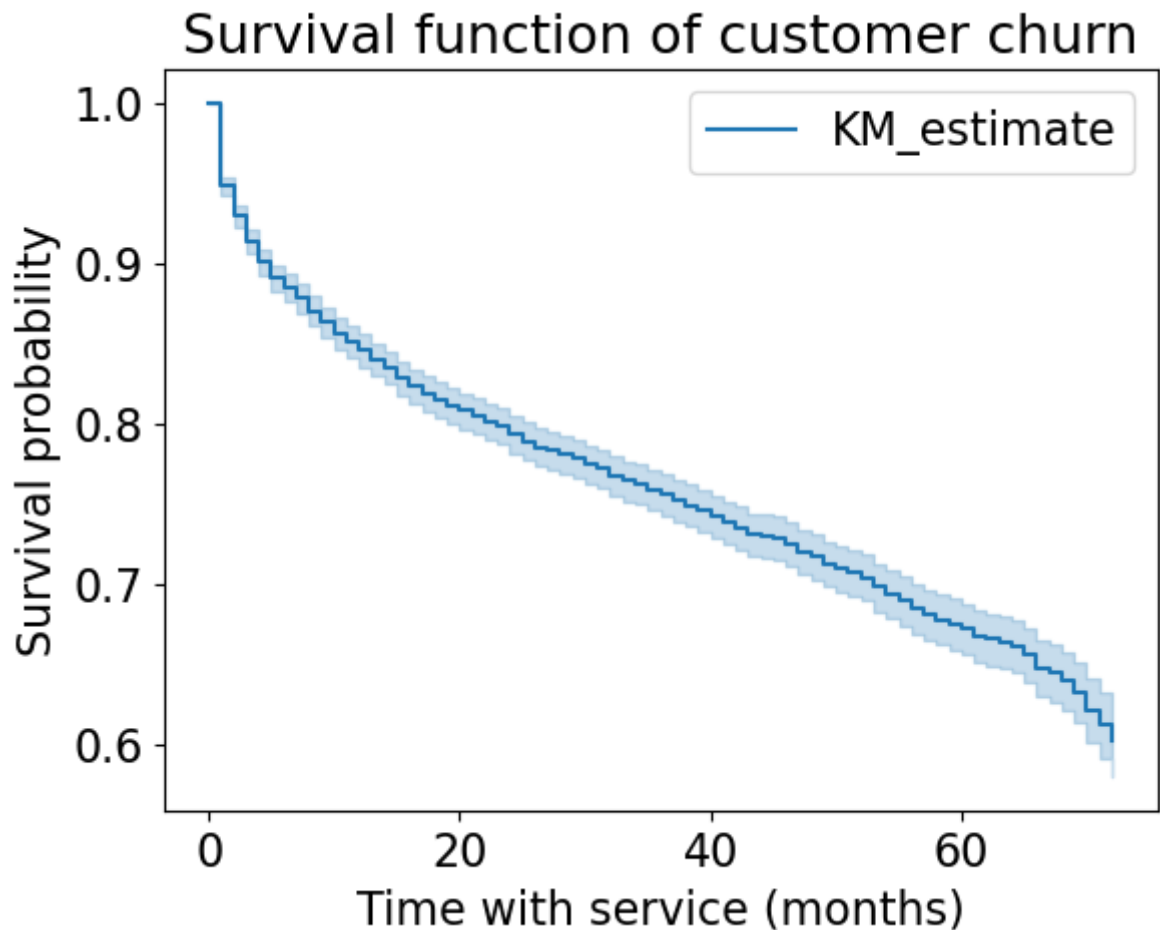
```
In [68]: 1 train_df[y_train == "No"]["tenure"].hist(grid=False)
2 plt.xlabel("months");
```



- Since the data was collected at a fixed time and these are the people who hadn't yet churned, those with larger `tenure` values here must have joined earlier.

Lifelines can also give us some "error bars":

```
In [69]: 1 kmf.plot()  
2 plt.title("Survival function of customer churn")  
3 plt.xlabel("Time with service (months)")  
4 plt.ylabel("Survival probability");
```




- We already have some actionable information here.
- The curve drops down fast at the beginning suggesting that people tend to leave early on.
- If there would have been a big drop in the curve, it means a bunch of people left at that time (e.g., after a 1-month free trial).
- BTW, the [original paper by Kaplan and Meier](https://web.stanford.edu/~lutan/coursepdf/KMpaper.pdf) (<https://web.stanford.edu/~lutan/coursepdf/KMpaper.pdf>) has been cited over 57000 times!

We can also create the K-M curve for different subgroups:

```
In [70]: 1 T = train_df_surv["tenure"]  
2 E = train_df_surv["Churn"]  
3 senior = train_df_surv["SeniorCitizen"] == 1
```



```
In [71]: 1 ax = plt.subplot(111)
2
3 kmf.fit(T[senior], event_observed=E[senior], label="Senior Citizens")
4 kmf.plot(ax=ax)
5
6 kmf.fit(T[~senior], event_observed=E[~senior], label="Non-Senior Citizens")
7 kmf.plot(ax=ax)
8
9 plt.ylim(0, 1)
10 plt.xlabel("Time with service (months)")
11 plt.ylabel("Survival probability");
```



- It looks like senior citizens churn more quickly than others.
- This is quite useful!

Cox proportional hazards model

- We haven't been incorporating other features in the model so far.
- The Cox proportional hazards model is a commonly used model that allows us to interpret how features influence a censored tenure/duration.

- You can think of it like linear regression for survival analysis: we will get a coefficient for each feature that tells us how it influences survival.
- It makes some strong assumptions (the proportional hazards assumption) that may not be true, but we won't go into this here.
- The proportional hazard model works multiplicatively, like linear regression with log-transformed targets.

```
In [72]: 1 cph = lifelines.CoxPHFitter()
2 cph.fit(train_df_surv, duration_col="tenure", event_col="Churn");

-----
--
LinAlgError                                Traceback (most recent call last)
File /opt/anaconda3/envs/cpsc330/lib/python3.10/site-packages/lifelines/fitters/coxph_fitter.py:1527, in SemiParametricPHFitter._newton_raphson_refit(self, X, T, E, weights, entries, initial_point, show_progress, step_size, precision, max_steps)
    1526 try:
-> 1527     inv_h_dot_g_T = spsolve(-h, g, assume_a="pos", check_finite=False)
1528 except (ValueError, LinAlgError) as e:

File /opt/anaconda3/envs/cpsc330/lib/python3.10/site-packages/scipy/linalg/_basic.py:254, in solve(a, b, sym_pos, lower, overwrite_a, overwrite_b, check_finite, assume_a, transposed)
    251 lu, x, info = posv(a1, b1, lower=lower,
    252                    overwrite_a=overwrite_a,
    253                    overwrite_b=overwrite_b)
    254 if info > 0:
    255     raise ValueError('Singular matrix: %s' % a)
```

- Ok, going that [that URL \(https://lifelines.readthedocs.io/en/latest/Examples.html#problems-with-convergence-in-the-cox-proportional-hazard-model\)](https://lifelines.readthedocs.io/en/latest/Examples.html#problems-with-convergence-in-the-cox-proportional-hazard-model), it seems the easiest solution is to add a penalizer.
 - FYI this is related to switching from LinearRegression to Ridge.
 - Adding drop='first' on our OHE might have helped with this.
 - (For 340 folks: we're adding regularization; lifelines adds both L1 and L2 regularization, aka elastic net)

```
In [73]: 1 cph = lifelines.CoxPHFitter(penalizer=0.1)
2 cph.fit(train_df_surv, duration_col="tenure", event_col="Churn");
```

We can look at the coefficients learned by the model and start interpreting them!

```
In [74]: 1 cph_params = pd.DataFrame(cph.params_).sort_values(by="coef", ascending
2 cph_params
```

	coef
covariate	
Contract_Month-to-month	0.812875
OnlineSecurity_No	0.311151
OnlineBackup_No	0.298561
PaymentMethod_Electronic check	0.280801
Partner_No	0.244814
...	...
OnlineBackup_Yes	-0.282600
PaymentMethod_Credit card (automatic)	-0.302801
OnlineSecurity_Yes	-0.330346
Contract_One year	-0.351821
Contract_Two year	-0.776427

```
In [78]: 1 X_train.drop(columns=["tenure"]).head()
```

```
Out[78]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	PhoneService	MultipleLines	Internet
6464	4726-DLWQN	Male	1	No	No	Yes	Yes	
5707	4537-DKTAL	Female	0	No	No	Yes	No	
3442	0468-YRPXN	Male	0	No	No	Yes	No	Fit
3932	1304-NECVQ	Female	1	No	No	Yes	Yes	Fit
6124	7153-CHRBV	Female	0	Yes	Yes	Yes	No	

I'm redefining feature types and our preprocessor for our sanity.

```
In [81]: 1 numeric_features = ["MonthlyCharges", "TotalCharges"]
2 drop_features = ["customerID", "tenure"]
3 passthrough_features = ["SeniorCitizen"]
4 target_column = ["Churn"]
5 # the rest are categorical
6 categorical_features = list(
7     set(train_df.columns)
8     - set(numeric_features)
9     - set(passthrough_features)
10    - set(drop_features)
11    - set(target_column)
12 )
```

```
In [82]: 1 preprocessor = make_column_transformer(
2     (
3         make_pipeline(SimpleImputer(strategy="median"), StandardScaler(
4             numeric_features,
5         )),
6         (OneHotEncoder(handle_unknown="ignore"), categorical_features),
7         ("passthrough", passthrough_features),
8         ("drop", drop_features),
9     )
10 )
```

```
In [83]: 1 preprocessor.fit(X_train);
```

```
In [86]: 1 new_columns = (
2     numeric_features
3     + preprocessor.named_transformers_["onehotencoder"]
4     .get_feature_names_out(categorical_features)
5     .tolist()
6     + passthrough_features
7 )
```

```
In [87]: 1 lr = make_pipeline(preprocessor, LogisticRegression(max_iter=1000))
2 lr.fit(X_train, y_train)
3 lr_coefs = pd.DataFrame(
4     data=np.squeeze(lr[1].coef_), index=new_columns, columns=["Coefficient"]
5 )
```

```
In [88]: 1 lr_coefs.sort_values(by="Coefficient", ascending=False)
```

```
Out[88]:
```


	Coefficient
Contract_Month-to-month	0.787653
InternetService_Fiber optic	0.600509
OnlineSecurity_No	0.291008
StreamingTV_Yes	0.258659
PaymentMethod_Electronic check	0.251646
...	...
MultipleLines_No	-0.169654
PaymentMethod_Credit card (automatic)	-0.204406
InternetService_DSL	-0.461593
TotalCharges	-0.743315
Contract_Two year	-0.765519

44 rows × 1 columns

- There is some agreement, which is good.
- But our survival model is much more useful.
 - Not to mention more correct.

- One thing we get with `lifelines` is confidence intervals on the coefficients:

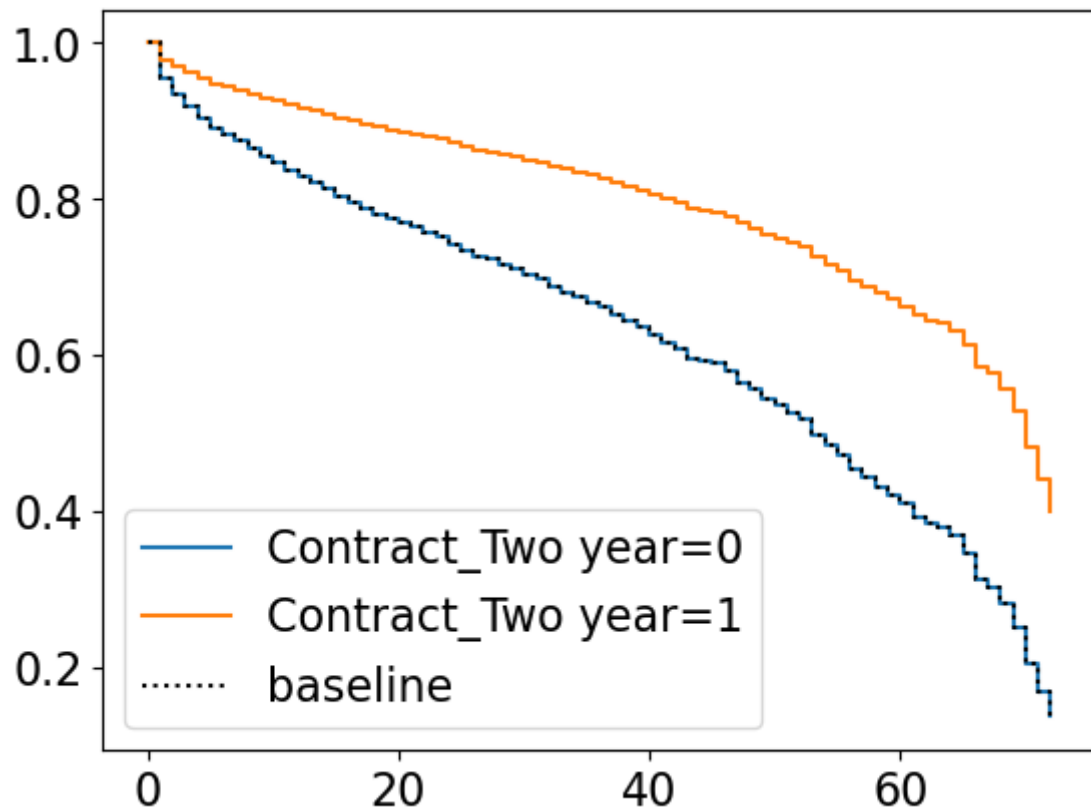
```
In [89]: 1 plt.figure(figsize=(8, 8))
2 cph.plot();
```



- (We could probably get the same for logistic regression if using `statsmodels` instead of `sklearn`.)
- However, in general, I would be careful with all of this.
- Ideally we would have more statistical training when using `lifelines` - there is a lot that can go wrong.
 - It comes with various diagnostics as well.
- But I think it's very useful to know about survival analysis and the availability of software to deal with it.
- Oh, and there are lots of other nice plots.

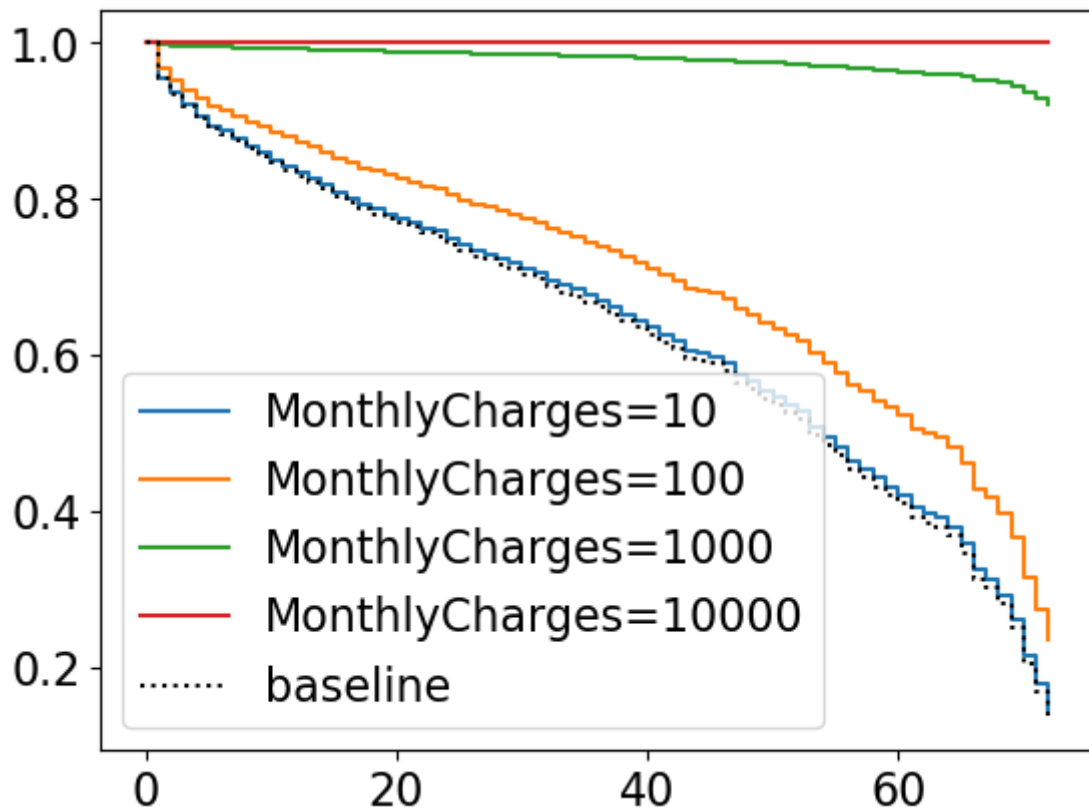
- Let's look at the survival plots for the people with
 - two-year contract (`Contract_Two year = 1`) and
 - people without two-year contract (`Contract_Two year = 0`)
- As expected, the former survive longer.

```
In [90]: 1 cph.plot_partial_effects_on_outcome("Contract_Two year", [0, 1]);
```



Now let's look at the survival plots for the people with different MonthlyCharges.

```
In [91]: 1 cph.plot_partial_effects_on_outcome("MonthlyCharges", [10, 100, 1000, 10000])
```



- That's the thing with linear models, they can't stop the growth.
- We have a negative coefficient associated with `MonthlyCharges`

```
In [92]: 1 cph_params.loc["MonthlyCharges"]
```

```
Out[92]: coef    -0.003185  
          Name: MonthlyCharges, dtype: float64
```

If your monthly charges are huge, it takes this to the extreme and thinks you'll basically never churn.

Prediction

- We can use survival analysis to make predictions as well.
- Here is the expected number of months to churn for the first 5 customers in the test set:


```
In [93]: 1 test_df_surv.drop(columns=["tenure", "Churn"]).head()
```

```
Out[93]:
```

	SeniorCitizen	MonthlyCharges	TechSupport_No	TechSupport_No internet service	TechSupport_Yes	MultipleLir
941	0.0	-1.154900	1.0	0.0	0.0	
1404	0.0	-1.383246	0.0	1.0	0.0	
5515	0.0	-1.514920	0.0	1.0	0.0	
3684	0.0	0.351852	1.0	0.0	0.0	
7017	0.0	-1.471584	0.0	1.0	0.0	

5 rows × 43 columns

```
In [94]: 1 test_df_surv.head()
```

```
Out[94]:
```

	Churn	tenure	SeniorCitizen	MonthlyCharges	TechSupport_No	TechSupport_No internet service	TechSupport
941	0.0	13.0	0.0	-1.154900	1.0	0.0	
1404	0.0	35.0	0.0	-1.383246	0.0	1.0	
5515	0.0	18.0	0.0	-1.514920	0.0	1.0	
3684	0.0	43.0	0.0	0.351852	1.0	0.0	
7017	0.0	51.0	0.0	-1.471584	0.0	1.0	

5 rows × 45 columns


How long each non-churned customer is likely to stay according to the model assuming that they just joined right now?

```
In [95]: 1 cph.predict_expectation(test_df_surv).head() # assumes they just joined
```

```
Out[95]: 941      35.206724
1404      69.023086
5515      68.608565
3684      27.565062
7017      67.890933
dtype: float64
```

Survival curves for first 5 customers in the test set:

```
In [96]: 1 cph.predict_survival_function(test_df_surv[:5]).plot()  
2 plt.xlabel("Time with service (months)")  
3 plt.ylabel("Survival probability");
```



From `predict_survival_function` documentation:

Predict the survival function for individuals, given their covariates. This assumes that the individual just entered the study (that is, we do not condition on how long they have already lived for.)

So these curves are "starting now".

- There's no probability prerequisite for this course, so this is optional material.
- But you can do some interesting stuff here with conditional probabilities.
- "Given that a customer has been here 5 months, what's the outlook?"
 - It will be different than for a new customer.
 - Thus, we might still want to predict for the non-churned customers in the training set!
 - Not something we really thought about with our traditional supervised learning.

Let's get the customers who have not churned yet.

```
In [97]: 1 train_df_surv_not_churned = train_df_surv[train_df_surv["Churn"] == 0]
```

We can *condition* on the person having been around for 20 months.

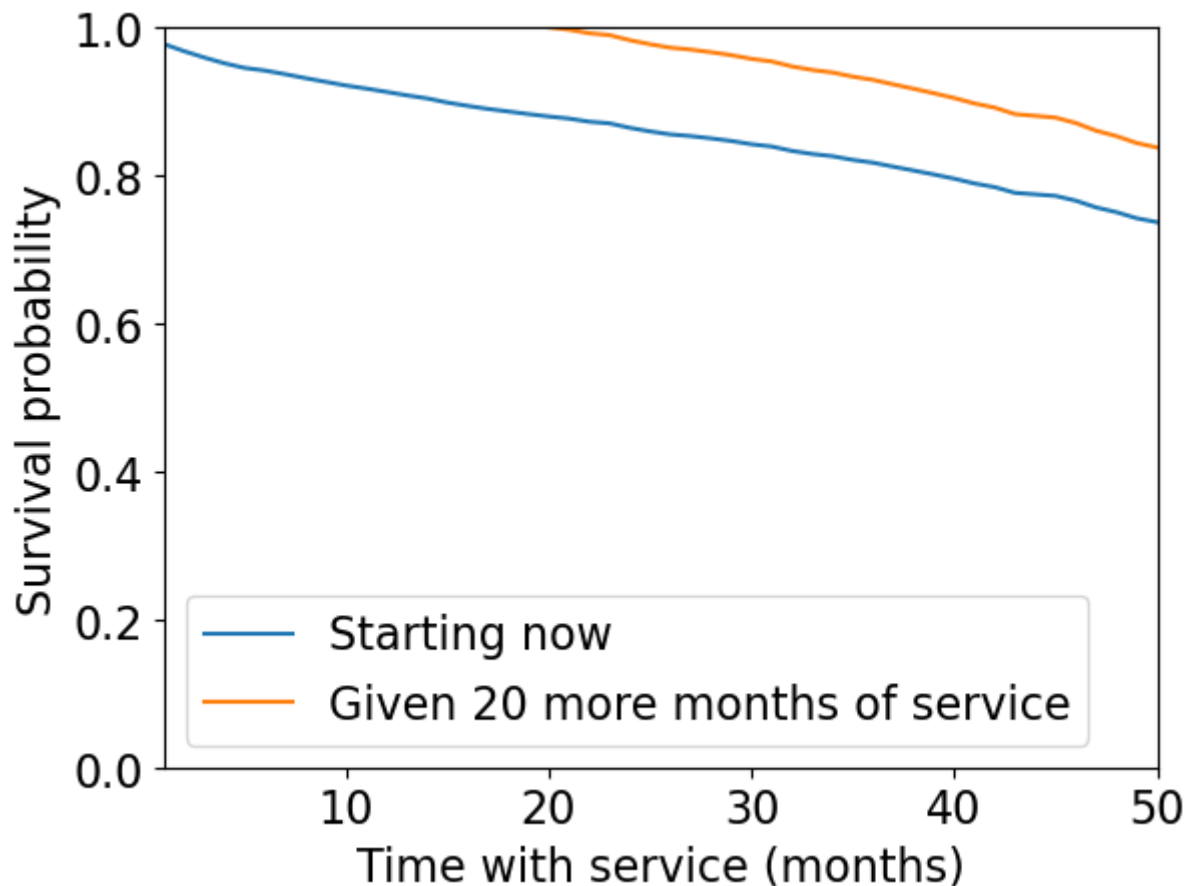
```
In [98]: 1 cph.predict_survival_function(train_df_surv_not_churned[:1], conditiona
```

Out[98]:

	6464
0.0	1.000000
1.0	0.996788
2.0	0.991966
3.0	0.989443
4.0	0.982570
...	...
68.0	0.429634
69.0	0.429634
70.0	0.429634
71.0	0.429634
72.0	0.429634

73 rows × 1 columns


```
In [99]: 1 plt.figure()
2 cph.predict_survival_function(train_df_surv_not_churned[:1]).plot(ax=pl
3 preds = cph.predict_survival_function(
4     train_df_surv_not_churned[:1], conditional_after=20
5 )
6 plt.plot(preds.index[20:], preds.values[:20])
7 plt.xlabel("Time with service (months)")
8 plt.ylabel("Survival probability")
9 plt.legend(["Starting now", "Given 20 more months of service"])
10 plt.ylim([0, 1])
11 plt.xlim([1, 50]);
```



- Look at how the survival function (and expected lifetime) is much longer *given* that the customer has already lasted 20 months.

- How long each non-churned customer is likely to stay according to the model assuming that they have been here for the tenure time?
- So, we can set this to their actual tenure so far to get a prediction of what will happen going forward:

```
In [100]: 1 cph.predict_survival_function(  
2     train_df_surv_not_churned[:1],  
3     conditional_after=train_df_surv_not_churned[:1]["tenure"],  
4 ).plot()  
5 plt.xlabel("Time into the future (months)")  
6 plt.ylabel("Survival probability")  
7 plt.ylim([0, 1])  
8 plt.xlim([0, 20]);
```



- Another useful application: you could ask what is the [customer lifetime value](https://en.wikipedia.org/wiki/Customer_lifetime_value) (https://en.wikipedia.org/wiki/Customer_lifetime_value).
 - Basically, how much money do you expect to make off this customer between now and when they churn?
- With regular supervised learning, tenure was a feature and we could only predict whether or not they had churned by then.

```
In [ ]: 1
```

Evaluation

By default score returns "partial log likelihood":

```
In [101]: 1 cph.score(train_df_surv)
```

```
Out[101]: -1.8641864337292489
```

```
In [102]: 1 cph.score(test_df_surv)
```

```
Out[102]: -1.7277854625841886
```

We can look at the "concordance index" which is more interpretable:

```
In [103]: 1 cph.concordance_index_
```

```
Out[103]: 0.8625888648969532
```

```
In [104]: 1 cph.score(train_df_surv, scoring_method="concordance_index")
```

```
Out[104]: 0.8625888648969532
```

```
In [105]: 1 cph.score(test_df_surv, scoring_method="concordance_index")
```

```
Out[105]: 0.8546143543902771
```

From the documentation [here](https://lifelines.readthedocs.io/en/latest/Survival%20Regression.html#model-selection-and-calibration-in-survival-regression)

(<https://lifelines.readthedocs.io/en/latest/Survival%20Regression.html#model-selection-and-calibration-in-survival-regression>):

Another censoring-sensitive measure is the concordance-index, also known as the c-index. This measure evaluates the accuracy of the ranking of predicted time. It is in fact a generalization of AUC, another common loss function, and is interpreted similarly:

- 0.5 is the expected result from random predictions,
- 1.0 is perfect concordance and,
- 0.0 is perfect anti-concordance (multiply predictions with -1 to get 1.0)

[Here \(https://stats.stackexchange.com/a/478305/11867\)](https://stats.stackexchange.com/a/478305/11867) is an excellent introduction & description of the c-index for new users.

```
In [96]: 1 # cph.log_likelihood_ratio_test()
```

```
In [97]: 1 # cph.check_assumptions(df_train_surv)
```

Other approaches / what did we not cover?

There are many other approaches to modelling in survival analysis:

- Time-varying proportional hazards.
 - What if some of the features change over time, e.g. plan type, number of lines, etc.
- Approaches based on deep learning, e.g. the [pysurvival](https://square.github.io/pysurvival/) (<https://square.github.io/pysurvival/>) package.
- Random survival forests.
- And more...

Types of censoring

There are also various types and sub-types of censoring we didn't cover:

- What we saw today is data with "right censoring"
- Sub-types within right censoring
 - Did everyone join at the same time?
 - Other there other reasons the data might be censored at random times, e.g. the person died?
- Left censoring
- Interval censoring

Summary

- Censoring and incorrect approaches to handling it
 - Throw away people who haven't churned
 - Assume everyone churns today
- Predicting tenure vs. churned
- Survival analysis encompasses both of these, and deals with censoring
- And it can make rich and interesting predictions!
- KM model -> doesn't look at features
- CPH model -> like linear regression, does look at the features

True/False questions

1. If all customers joined a service at the same time (hypothetically), then censoring would not be an issue.
2. The Cox proportional hazards model (cph above) assumes the effect of a feature is the same for all customers and over all time.
3. Survival analysis can be useful even without a "deployment" stage.

References

Some people working with this same dataset:

- <https://medium.com/@zachary.james.angell/applying-survival-analysis-to-customer-churn-40b5a809b05a> (<https://medium.com/@zachary.james.angell/applying-survival-analysis-to-customer-churn-40b5a809b05a>)
- <https://towardsdatascience.com/churn-prediction-and-prevention-in-python-2d454e5fd9a5> (<https://towardsdatascience.com/churn-prediction-and-prevention-in-python-2d454e5fd9a5>) (Cox)
- <https://towardsdatascience.com/survival-analysis-in-python-a-model-for-customer-churn-e737c5242822> (<https://towardsdatascience.com/survival-analysis-in-python-a-model-for-customer-churn-e737c5242822>)
- <https://towardsdatascience.com/survival-analysis-intuition-implementation-in-python-504fde4fcf8e> (<https://towardsdatascience.com/survival-analysis-intuition-implementation-in-python-504fde4fcf8e>)

lifelines documentation:

- <https://lifelines.readthedocs.io/en/latest/Survival%20analysis%20with%20lifelines.html> (<https://lifelines.readthedocs.io/en/latest/Survival%20analysis%20with%20lifelines.html>)
- <https://lifelines.readthedocs.io/en/latest/Survival%20Analysis%20intro.html#introduction-to-survival-analysis> (<https://lifelines.readthedocs.io/en/latest/Survival%20Analysis%20intro.html#introduction-to-survival-analysis>)